

## The Problem we are trying to solve

Solve for  $\tilde{x}$ :

$$A(x)\tilde{x} = b(x) \quad (1)$$

with  $A(x) = (B^{-1} + (HM)^T R^{-1} (HM))$

- Learn a low-rank approximation using DNN (to use as preconditioner):
  - $x \mapsto (U(x), S(x)) \in \mathbb{R}^{n \times r} \times \mathbb{R}^r$  where  $A(x) \approx U(x)\text{diag}(S(x))U(x)^T$
  - Approximate the norm using random vectors
  - Look for an approximation in the dual space instead ?
- Better architecture for neural network (CNN, UNet, attention layers?), instead of MLP
- (Technical stuff): DVC and MLflow for reproducibility, data versioning and workflow management

# Variational DA and ML

Using ML-based preconditioners in VarDA problems

---

Victor Trappler

April 3, 2023



# Table of contents

1. Introduction
2. Variational Data Assimilation
3. Data-driven preconditioning
4. Neural Network architecture
5. Example on Lorenz96 (40D)

# Introduction

---

# Notation and setting

- $x \in \mathbb{R}^n$  state vector: Variables that describe a physical system ( $n = \mathcal{O}(10^{6-9})$ )

Different sources of information on a state vector  $x$ :


- A priori information  $p(x)$ 
  - Historical data
  - Balance equations
- Observations  $y$ , obtained more or less indirectly
- Numerical model which maps the state to the observations  $\mathcal{G} = \mathcal{H} \circ \mathcal{M}$

How to combine them ? Bayes theorem

# Data Assimilation and Bayesian Inference

Bayes' theorem : update information on  $x$  using  $y$

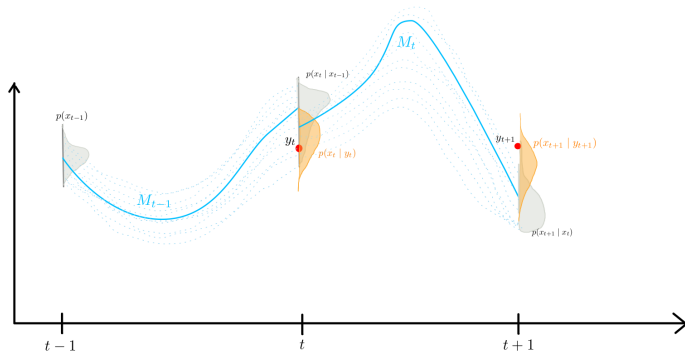
$$p(x | y) = p(x) \frac{\overbrace{p(y | x)}^{\text{likelihood}}}{\underbrace{p(y)}_{\text{evidence}}} \quad (2)$$



# Data Assimilation and Bayesian Inference

Bayes' theorem sequentially: update information on  $x_t$  using  $y_t$

$$p(x_t | y_t) = \underbrace{p(x_t | x_{t-1})}_{\text{prior/forecast}} \underbrace{\frac{p(y_t | x_t)}{p(y_t)}}_{\text{evidence}} \quad (2)$$



# Variational Data Assimilation

---



We are interested in a point estimate of the posterior distribution: Maximum A Posteriori

$$\min_x \{-\log p(x | y)\} = \min_x \left\{ \underbrace{-\log p(y | x)}_{\text{log-lik=misfit}} - \underbrace{\log p(x)}_{\text{regularization}} \right\} \quad (3)$$

- $Y | x$  encodes the relation between the forward model and the observations
- $X$  encodes the knowledge we gathered so far on  $x$

## Gaussian Assumptions

- $y = \mathcal{G}(x) + \varepsilon$ , then  $Y | x \sim \mathcal{N}(\mathcal{G}(x), R)$
- $x \sim \mathcal{N}(x^b, B)$

# Standard formulation of the objective function

Using the Gaussian assumptions, we have  $p(x | y) \propto e^{-J(x)}$  with

$$J(x) = \frac{1}{2} \|\mathcal{G}(x) - y\|_{R^{-1}}^2 + \frac{1}{2} \|x - x^b\|_{B^{-1}}^2 \quad (4)$$

Which can be simplified to

$$J(x) = \frac{1}{2} \|\mathcal{G}(x) - y\|^2 \quad (5)$$

## Analysis

The analysis is the MAP (point estimate)

$$x^a = \min_x J(x) = \min_x \frac{1}{2} \|\mathcal{G}(x) - y\|^2 \quad (6)$$

which is a non-linear least square problem

→  $\mathcal{G}$  is a numerical model, expensive to evaluate. How to minimize  $J$  ?

## Incremental formulation

$$J_{\text{inc}}(\delta x; x) = J(x) + \underbrace{\nabla J^T \delta x + \frac{1}{2}(\delta x)^T H(\delta x)}_{\text{quadratic wrt } \delta x} \approx J(x + \delta x) \quad (7)$$

with

$$G = \nabla \mathcal{G} = \text{Tangent Linear of } \mathcal{G} \text{ at } x \quad (8)$$

$$\nabla J(x) = G^T(\mathcal{G}(x) - y) = G^T d \quad (9)$$

$$H(x) = G^T G + \underbrace{Q(x)}_{\frac{\partial^2}{\partial x^2}} \approx G^T G = H_{\text{GN}} \quad (10)$$

To minimize the quadratic approx., the increment  $\delta x$  is the solution to the linear system

$$H_{\text{GN}} \delta x = -\nabla J \iff (G^T G) \delta x = -G^T d \quad (11)$$

# Nested Loops

Set  $k = 0$


Repeat until convergence/computational budget spent

## Outer Loop

- Evaluate
  - Forward  $\mathcal{G}(x^{(k)})$
  - Tangent Linear  $G$
  - Objective  $J(x^{(k)})$
  - Gradient  $G^T d = G^T(\mathcal{G}(x^{(k)}) - y)$

## Inner Loop

- Solve  $(G^T G)\delta x^{(k)} = -G^T d$
- $x^{(k+1)} = x^{(k)} + \delta x^{(k)}$



# Nested Loops

Set  $k = 0$

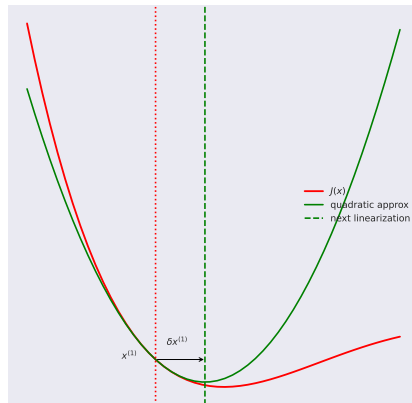
Repeat until convergence/computational budget spent

## Outer Loop

- Evaluate
  - Forward  $\mathcal{G}(x^{(k)})$
  - Tangent Linear  $G$
  - Objective  $J(x^{(k)})$
  - Gradient  $G^T d = G^T(\mathcal{G}(x^{(k)}) - y)$

## Inner Loop

- Solve  $(G^T G) \delta x^{(k)} = -G^T d$
- $x^{(k+1)} = x^{(k)} + \delta x^{(k)}$




# Nested Loops

Set  $k = 0$

Repeat until convergence/computational budget spent

## Outer Loop

- Evaluate
  - Forward  $\mathcal{G}(x^{(k)})$
  - Tangent Linear  $G$
  - Objective  $J(x^{(k)})$
  - Gradient  $G^T d = G^T(\mathcal{G}(x^{(k)}) - y)$
- Inner Loop
  - Solve  $(G^T G) \delta x^{(k)} = -G^T d$
  - $x^{(k+1)} = x^{(k)} + \delta x^{(k)}$




# Nested Loops

Set  $k = 0$

Repeat until convergence/computational budget spent

## Outer Loop

- Evaluate
  - Forward  $\mathcal{G}(x^{(k)})$
  - Tangent Linear  $G$
  - Objective  $J(x^{(k)})$
  - Gradient  $G^T d = G^T(\mathcal{G}(x^{(k)}) - y)$
- Inner Loop
  - Solve  $(G^T G) \delta x^{(k)} = -G^T d$
  - $x^{(k+1)} = x^{(k)} + \delta x^{(k)}$



- Trying to learn the forward operator  $\mathcal{G} = \mathcal{H} \circ \mathcal{M}$  is not worth the effort
  - Very costly, complex, high-dimensional
  - $d = \mathcal{G}(x) - y$  is central to compute the gradient (ie the direction of descent)
- Use ML to speed up inner loop ?
  - Need to solve  $(G^T G) \delta x = -G^T d$
  - How to speed-up iterative methods ?
    - Reduce the number of iterations to convergence
    - Reduce error for constant number of iterations
- Dimension reduction
  - Find a lower dimensional representation of the state space  $x \in \mathcal{X}$
  - Lower-dimensional Manifold on which the optimization can take place
  - Non-linear dimension reduction with diffeomorphism



## Data-driven preconditioning

---

# Solving linear systems

The increment  $\delta x$  verifies

$$\underbrace{(G^T G)}_A \delta x = \underbrace{-G^T d}_b \quad (12)$$

and using Conjugate Gradient, the error at the  $k$ th iteration is bounded according to

$$\|e_k\| \leq 2 \left( \frac{\sqrt{\kappa(A)} - 1}{\sqrt{\kappa(A)} + 1} \right)^k \|e_0\| \quad (13)$$

where  $\kappa(A) = \left| \frac{\text{largest eigenvalue}}{\text{smallest eigenvalue}} \right| \geq \kappa(I_n) = 1$

$\Rightarrow$  Smaller  $\kappa$  = better rate of convergence

(More generally, depends on the whole distribution of the eigenvalues)

[Haben et al., 2011, Gürol et al., 2014, Tabeart et al., 2021, Robert et al., 2006]

# Desired properties of preconditioners

Let  $\delta x$  be a solution of

$$A\delta x = b \quad (14)$$

## Left Preconditioner

Let  $H^{-1}$  be an invertible matrix

$\delta x$  is also a solution of

$$(H^{-1}A)\delta x = H^{-1}b \quad (15)$$

and we hope that the new linear system is easier to solve

Desired properties:

- $H^{-1}$  symmetric and invertible
- $H^{-1}A$  should be close to  $I_n$
- $\kappa(H^{-1}A) < \kappa(A)$

# State dependent preconditioner

One-size-fits-all preconditioners do not exist (or are very simplistic).

Recalling that  $A = G^T G = G(x)^T G(x)$  is state-dependent ( $G$  TL of the forward model)

## State-dependent preconditioner

$$x \mapsto \text{prec}(x) = H^{-1}(x) \quad (16)$$

which exploits the fact that  $G(x)^T G(x)$  is not arbitrary

- Defined according to a numerical model
- Positive-definite matrix and symmetric

We want  $H^{-1}A$  close to  $I_n$

What loss to choose ?

- $\|H^{-1} - A^{-1}\|$   
→  $A^{-1}$  is what we are trying to avoid computing...
- $\|I_n - H^{-1}A\|$   
→ quite "unstable" objective in practice, especially for badly conditioned system (local minimum is  $H^{-1} = 0$ )
- $\|H - A\|$   
→ but we need to train using  $H$ , and use  $H^{-1}$  as a preconditioner

# Neural Network architecture


---

# Low-rank updates

Let  $S = (s_1, \dots, s_r)$  be  $r$  vectors of  $\mathbb{R}^n$ ,  $r \ll n$

$$H_{\text{LR}}(S) = I_n + SS^T \quad (17)$$

$$H_{\text{LR}}^{-1}(S) = I_n - \underbrace{S(I_r - S^T S)^{-1} S^T}_{\text{inv of dim } r} \quad (18)$$



**Figure 1:** Flowchart for prec. training using low rank matrices

→ Did not give results (for a lack of structure ?)

# Deflation-like preconditioner

$S = (s_1, \dots, s_r)$  has  $r$  orthonormal columns ( $S^T S = I_r$ ), and  $(\lambda_1, \dots, \lambda_r) \in \mathbb{R}^r$ ,  $\lambda_i > 0$

$$H_{\text{defl}}(S, \lambda) = I_n + \sum_{i=1}^r (\lambda_i - 1) s_i s_i^T \quad (19)$$

$$H_{\text{defl}}^{-1}(S, \lambda) = I_n + \sum_{i=1}^r \left( \frac{1}{\lambda_i} - 1 \right) s_i s_i^T \quad (20)$$

$$(21)$$




Figure 2: Flowchart for prec. training using deflation



# Limited Memory Preconditioners [Tshimanga, 2007, Gratton et al., 2011]

Let  $S = (s_1, \dots, s_r)$  be  $r$  vectors of  $\mathbb{R}^n$ ,  $r \ll n$

$$H_{\text{LMP}}^{-1}(S, \text{AS}) = \left( \begin{bmatrix} I_n \end{bmatrix} - \begin{bmatrix} S \end{bmatrix} \begin{bmatrix} \Sigma \end{bmatrix} \begin{bmatrix} \text{AS} \end{bmatrix}^T \right) \left( \begin{bmatrix} I_n \end{bmatrix} - \begin{bmatrix} \text{AS} \end{bmatrix} \begin{bmatrix} \Sigma \end{bmatrix} \begin{bmatrix} S \end{bmatrix}^T \right) + \mu \begin{bmatrix} S \end{bmatrix} \begin{bmatrix} \Sigma \end{bmatrix} \begin{bmatrix} S \end{bmatrix}^T \quad (22)$$

$$H_{\text{LMP}}(S, \text{AS}) = \begin{bmatrix} I_n \end{bmatrix} + \begin{bmatrix} \text{AS} \end{bmatrix} \begin{bmatrix} \Sigma \end{bmatrix} \begin{bmatrix} \text{AS} \end{bmatrix}^T - \frac{1}{\mu} \begin{bmatrix} S \end{bmatrix} \begin{bmatrix} \Gamma \end{bmatrix} \begin{bmatrix} S \end{bmatrix}^T \quad (23)$$

Given  $S \in \mathbb{R}^{n \times r}$  and  $\text{AS} \in \mathbb{R}^{n \times r}$ , we can construct  $H_{\text{LMP}}$  and  $H_{\text{LMP}}^{-1}$  directly, with nice properties (even though  $\Sigma$  and  $\Gamma$  require the inversion of a  $r \times r$  matrix)

Flexible preconditioners:

- $(s_1, \dots, s_r)$  eigenvalues of  $A \rightarrow$  Spectral Preconditioner
- $(s_1, \dots, s_r)$  CG descent vectors  $\rightarrow$  QN preconditioner

# LMP construction using leap-forward NN




Figure 3: Neural Network architecture for LMP

**But:** No particular structure imposed for  $S' \Rightarrow (H_{\text{LMP}}(S, S'))^{-1} \neq H_{\text{LMP}}^{-1}(S, S')$

However: Good results were still obtained, even though the exact inverse is not used

# LMP: self adjoint variation

Force the self-adjointness of the operator  $S \mapsto S'$  (which should be  $= AS$ ) by constructing at the same time a low-rank approximation of  $A$




Figure 4: Neural Network architecture for symmetric LMP

with  $L(x) \in \mathbb{R}^{n \times n'}$

## Training setting: Dataset

For an input  $x$ , we assume to have access to  $G(x)^T G(x)$ ,

→ Training dataset:  $\mathcal{D} = \{(\underbrace{x_i}_{\in \mathbb{R}^n}, \underbrace{G(x_i)^T G(x_i)}_{\in \mathbb{R}^{n \times n}})\}$

- Very large in memory when dimension increases
- We access  $G$  only as an operator:  $\text{TL}(x, z) = G(x) \cdot z$
- Same for adjoint:  $\text{Adj}(x, y) = G(x)^T \cdot y$
- Constructing  $G(x)$  would require  $n$  call to the TL

## Less storage intensive solution: Iterable Datasets

Estimate the  $L_2$  norm using random Gaussian vectors:

### Matrix norm estimation

For a matrix  $M$ , and  $z \sim \mathcal{N}(0, I)$

$$\mathbb{E}_Z [\|Mz\|_2^2] = \|M\|_F^2 \quad (24)$$

→ Iterable Dataset:  $\mathcal{D} = \{(\underbrace{x_i}_{\in \mathbb{R}^n}, \underbrace{Z_i}_{\in \mathbb{R}^{n \times n_z}}, \underbrace{G(x_i)^T G(x_i) Z_i}_{\in \mathbb{R}^{n \times n_z}})_i\}$  where  $Z_i$  has  $n_z$  columns iid  $\mathcal{N}(0, I)$

The loss for a data point becomes then

$$\mathcal{L}_\theta(x_i) = \sum_{j=1}^{n_z} \|H_\theta(x_i) z_i^j - G^T(x_i) G(x_i) z_i^j\|_2^2 \quad (25)$$

and we can generate the dataset "on the fly", and train the network in an online manner

## A-conjugacy

Let  $S = (s_1 | s_2 | \dots | s_r)$ . The column vectors of  $S$  are A-conjugate when  $S^T A S$  is diagonal

Since  $S'$  is supposed to be  $AS$

$$K = S^T S' \text{ has general term } K_{i,j} = \langle s_i, s'_j \rangle = s_i^T s'_j \approx s_i^T A s_j \quad (26)$$

$$R_{\text{conjugacy}}(S) \propto \| \begin{array}{|c|} \hline K \\ \hline \end{array} \| \quad (27)$$

But we may also add regularization on symmetry of the LMP, orthonormalization of the outputs etc...

## Example on Lorenz96 (40D)

---

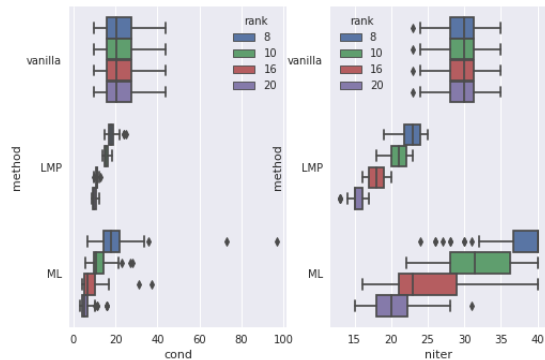



Figure 5: Statistics on inference model



# Inner loop CV: leap forward LMP ( $r = 20$ )

Comparison:


- **Baseline**: no preconditioner
- **Spectral LMP**: computation of eigenvectors for every inner loop
- **ML-LMP**: preconditioned inner loop using "leap forward" LMP



# Inner loop CV: leap forward LMP ( $r = 16$ )

Comparison:


- **Baseline**: no preconditioner
- **Spectral LMP**: computation of eigenvectors for every inner loop
- **ML-LMP**: preconditioned inner loop using "leap forward" LMP



# Inner loop CV: self adjoint LMP ( $r = 20$ )




Comparison:




- **Baseline**: no preconditioner
- **Spectral LMP**: computation of eigenvectors for every inner loop
- **ML-LMP**: preconditioned inner loop using self adjoint LMP




## Work in progress / Open questions

- Apply this preconditioner in 4D-Var ✓
- Output both  $S_\theta$  and  $AS_\theta$  ✓
- How to train without explicit access to  $A$  ⚙️
  - Online training ?
  - Use information of  $G^T(x)G(x)Z$  (Link to REVD [Daužickaitė et al., 2021])
- How to get more consistent results in 4D-Var ? ⚙️
- Which regularization to use ⚙️
- Apply to system with worse conditioning (QG) and/or higher dimension (QG / KS) ⚙️
- Exploit conditioning using the prior (" $(B + G^T G) \rightarrow (I + \tilde{G}^T \tilde{G})$ ") ?
- Extension to weak 4DVar ?
- Adapt to time-varying observation operators ?

-  Daužickaitė, I., Lawless, A. S., Scott, J. A., and van Leeuwen, P. J. (2021).  
**Randomised preconditioning for the forcing formulation of weak constraint 4D-Var.**  
*Quarterly Journal of the Royal Meteorological Society*, 147(740):3719–3734.
-  Gratton, S., Laloyaux, P., Sartenaer, A., and Tshimanga, J. (2011).  
**A reduced and limited-memory preconditioned approach for the 4D-Var data-assimilation problem: Reduced 4D-Var using the LMP Preconditioner.**  
*Quarterly Journal of the Royal Meteorological Society*, 137(655):452–466.
-  Gürol, S., Weaver, A. T., Moore, A. M., Piacentini, A., Arango, H. G., and Gratton, S. (2014).  
**B -preconditioned minimization algorithms for variational data assimilation with the dual formulation: B -preconditioned minimization algorithms.**  
*Quarterly Journal of the Royal Meteorological Society*, 140(679):539–556.

-  Haben, S., Lawless, A., and Nichols, N. (2011).  
**Conditioning and preconditioning of the variational data assimilation problem.**  
*Computers & Fluids*, 46(1):252–256.
-  Robert, C., Blayo, E., and Verron, J. (2006).  
**Reduced-order 4D-Var: A preconditioner for the Incremental 4D-Var data assimilation method: A PRECONDITIONER FOR THE 4D-VAR METHOD.**  
*Geophysical Research Letters*, 33(18):n/a–n/a.
-  Tabcart, J. M., Dance, S. L., Lawless, A. S., Nichols, N. K., and Waller, J. A. (2021).  
**New bounds on the condition number of the Hessian of the preconditioned variational data assimilation problem.**

-  Tshimanga, J. (2007).  
On a class of limited memory preconditioners for large-scale nonlinear  
least-squares problems (with application to variational ocean data assimilation).